Theoretical Architecture and Systems Analysis of the Tokenized Bio Hedge Fund and Associated Computational Ecosystems

Introduction to the Convergence of Decentralized Science and Predictive Processing

The contemporary landscape of computational finance, distributed ledger technology, and artificial intelligence is undergoing a profound structural convergence. Decentralized Finance (DeFi), multi-agent artificial intelligence networks, and advanced theoretical frameworks for predictive processing are merging to create highly autonomous, self-optimizing economic entities. This report provides an exhaustive architectural analysis of this convergence by examining the digital footprint, open-source repositories, and theoretical inclinations of the developer profile coad1024-cmd (also known as Dzuname). Although the specific repository titled "Tokenized-Bio-Hedge-Fund" remains inaccessible or currently unavailable to the public, the developer's surrounding body of work provides a robust foundation for reconstructing the likely systemic architecture of such a platform.

The developer's public ecosystem demonstrates a sophisticated command of production-grade multi-agent AI systems, advanced algorithmic tokenomics designed to mitigate severe market exploits, and a conceptual alignment with theoretical physics and cognitive science—most notably the Free Energy Principle (FEP) and Active Inference. By synthesizing the technical parameters of the developer's accessible repositories with the mathematical frameworks of Active Inference articulated by prominent researchers such as Karl Friston, this report projects the systemic design of a tokenized hedge fund specializing in biological and biotechnological assets. The analysis spans computational tokenomics, Model Context Protocol (MCP) integrations, subagent orchestration, causal incentive structures, and the rigorous application of Bayesian mechanics to financial modeling.

Through evaluating the repository structures alongside extensive symposium transcripts on Applied Active Inference, we can establish a comprehensive blueprint of how an autonomous, decentralized financial entity can leverage epistemic foraging, structure learning, and dynamic niche construction to navigate the high-uncertainty domain of biotechnology investments.

Baseline Ecosystem Analysis: The Developer Profile and Cyberphysical Orientation

The GitHub profile coad1024-cmd (Dzuname) reveals a specialized focus on autonomous systems and decentralized financial infrastructure. With 18 active repositories and a bio indicating a highly focused, perhaps asynchronous workflow ("I may be slow to respond"), the profile functions as a workspace for experimental and production-grade architectures. The

developer follows the causalincentives organization, placing their theoretical orientation squarely at the intersection of causal inference, AI alignment, and applied tokenomics. This alignment is critical, as causal inference moves beyond mere statistical correlation, seeking to model the underlying data-generating processes that govern complex systems—an approach that directly mirrors the generative models foundational to Active Inference.
The public repositories establish a clear trajectory toward building verifiable, autonomous, and market-resilient systems. Table 1 provides a structural overview of the most critical public repositories that inform the theoretical projection of the Tokenized Bio Hedge Fund.

Repository Name	Primary Language	Architectural Significance	Core Capabilities
TRAVEL-Planner	TypeScript	Multi-Agent Orchestration	13 specialized AI agents, 14 MCP tool servers, RAG knowledge base, Next.js frontend.
dual-penalty-[span_4](start_span)[span_4](end_span)bonding-curve	Python	Advanced Tokenomics	Production-ready Bonding Curve v1.1. Graduated penalties, market-aware exponential decay.
Stablecoin_Re[span_5](start_span)[span_5](end_span)search	Python	Financial Mechanism Design	Workspace dedicated to advanced stablecoin economic models and research.
skillproof-ba[span_6](start_span)[span_6](end_span)ckend	Python	Verifiable Asset Generation	AI Portfolio Factory CLI, simulated LLM logic, automated PDF generation via reportlab.

The intersection of these repositories points to a unified design philosophy. The deployment of specialized AI agents (orchestrated via MCP) serves to gather and process complex data, while rigorous cryptographic and economic mechanisms (bonding curves, stablecoins) are leveraged to manage the financial outcomes of that data processing. This multi-layered approach reflects a cyberphysical understanding of modern systems, where digital tools serve as natural extensions of collective cognitive processes.

Verifiable Asset Generation and Logic Sandboxing

To comprehend the operational execution layer of the hypothesized Tokenized Bio Hedge Fund, it is imperative to examine the skillproof-backend repository. This project is structured as an "AI Portfolio Factory" Command Line Interface (CLI) prototype designed to help candidates create verifiable portfolio assets based on real-time job market data. The architecture relies on Python 3.12+ and utilizes a highly modular design pattern, separating logic into data, ai, and utils directories to mimic production-ready structures.
The core workflow of this system operates in three distinct phases: Scan, Sprint, and Asset. During the Scan phase, the system analyzes simulated job descriptions to identify high-demand skills for a target role. In the Sprint phase, it generates a specific, time-bound challenge based

on those identified skills. Finally, in the Asset phase, it accepts a user's solution summary and utilizes the reportlab library to programmatically generate a professional PDF case study.
When extrapolating this workflow to a bio-focused hedge fund, the underlying pipeline transitions from human resource portfolio generation to rigorous financial due diligence. The "Scan" phase maps to the continuous ingestion and analysis of biomedical literature, FDA trial databases, and synthetic biology patent filings. The "Sprint" phase maps to targeted epistemic foraging, where the system identifies critical knowledge gaps in a biotech startup's clinical trial methodology and generates specific analytical queries to resolve that uncertainty. The "Asset" phase maps to the automated generation of an immutable, cryptographically verifiable "Proof of Due Diligence" report, which is deployed alongside capital allocation to justify the investment to the decentralized community of token holders. This architecture demonstrates the developer's capacity to build deterministic, verifiable outputs from non-deterministic AI inputs.

Advanced Computational Tokenomics: The Dual-Penalty Bonding Curve

A critical component in understanding the financial mechanics of the Tokenized Bio Hedge Fund is the developer's dual-penalty-bonding-curve repository. Automated Market Makers (AMMs) and bonding curves are fundamental primitives in decentralized finance, deployed via smart contracts to mint and burn tokens according to a predefined mathematical formula. These formulas dynamically adjust the asset price based on the circulating token supply, providing continuous liquidity without the need for a traditional order book. However, standard bonding curves—such as those following simple linear or exponential functions—are highly susceptible to adversarial exploitation, including front-running, rug pulls, and flash loan attacks.

Mitigating Systemic Vulnerabilities

Flash loan attacks in DeFi occur when a malicious actor borrows a massive amount of capital uncollateralized, manipulates the pricing curve of an AMM by executing massive buy and sell orders within a single transaction block, and extracts the resulting arbitrage value before repaying the loan. Standard bonding curves react purely to token supply changes without temporal awareness, making them ideal targets for such zero-duration exploits.
The dual-penalty-bonding-curve implemented by coad1024-cmd introduces a "v1.1" production-ready architecture featuring individual graduated penalties and market-aware exponential decay. This design explicitly neutralizes flash loans and rug pulls by introducing profound temporal and volume-based friction into the trading execution logic.
If we define the baseline price P as a function of token supply S, a standard polynomial bonding curve follows the logic:
P(S) = m \cdot S^n
The developer's innovation injects a complex penalty function, Pen(\Delta t, \Delta V), which dynamically scales the exit (sell) fee based on two critical dimensions: the time elapsed since the asset was initially purchased (\Delta t) and the percentage of the total liquidity pool being liquidated by the transaction (\Delta V).

Market-Aware Exponential Decay

The "market-aware exponential decay" mechanism is particularly sophisticated. It ensures that the penalty fee does not drop linearly—which could still be gamed by sophisticated algorithms—but decays exponentially over time, forcing long-term alignment from capital allocators. Furthermore, the decay rate adjusts dynamically based on broader market volatility, integrating exogenous market awareness into an endogenous pricing function.

$Fee(\Delta t) = F_{max} \cdot e^{-\lambda(\sigma) \cdot \Delta t}$

In this formulation, F_{max} represents the absolute maximum penalty fee imposed on immediate, high-volume liquidations, Δt is the holding duration of the specific wallet address, and $\lambda(\sigma)$ is a variable decay constant that is actively modulated by current market volatility σ.

In the context of a Tokenized Bio Hedge Fund, this tokenomic structure is paramount for survival. Biological assets, such as intellectual property for novel pharmaceuticals or longitudinal clinical trial data, are inherently illiquid and require extended capital horizons to reach maturity. A hedge fund tokenizing these complex assets must prevent short-term speculative attacks from draining the treasury before the biological assets yield returns. By deploying a dual-penalty bonding curve, the fund structurally ensures that capital entering the system is incentivized to remain through the multi-year duration of clinical trials or biotech research phases. It explicitly punishes short-term capital flight and mathematically prevents flash loan manipulation, thereby stabilizing the financial environment required for long-term scientific investment.

Table 2 highlights the structural differences between traditional bonding curves and the dual-penalty variant.

Feature	Standard Polynomial Curve	Dual-Penalty Bonding Curve v1.1
Price Determinant	Supply magnitude exclusively	Supply magnitude, time held, transaction volume, market volatility
Flash Loan Vulnerability	Extremely High (zero-duration extraction)	Eliminated (maximum penalty applied to zero-duration sales)
Capital Alignment	Short-term speculation incentivized	Long-term holding structurally mandated via exponential decay
Volatility Response	Static (curve shape is immutable)	Dynamic (λ decay constant adjusts to exogenous volatility)

Base Currency Stabilization and Risk Amortization

The developer's explicit focus on Stablecoin_Research provides an essential complementary layer to the bonding curve architecture. A hedge fund cannot effectively assess the value of volatile exogenous assets (biotech IP) if its own base unit of account is simultaneously volatile. The introduction of a stablecoin serves to amortize risk and provide a stable numerical foundation for the fund's internal calculus.

By researching and presumably developing bespoke stablecoin architectures, the developer ensures that the Tokenized Bio Hedge Fund possesses a reliable numéraire. This allows the system's analytical agents to compute exact valuations without the noise generated by underlying cryptocurrency fluctuations. In control theory, mitigating expected volatility in this manner is mathematically related to risk-sensitive control in economics, ensuring that the system's core operative baseline remains insulated from external financial shocks.

Multi-Agent Orchestration via the Model Context Protocol

The second foundational pillar of the developer's architecture is the deployment of vast, orchestrated multi-agent AI systems, as evidenced by the TRAVEL-Planner repository. This advanced system utilizes 13 specialized AI agents and 14 Model Context Protocol (MCP) tool servers, supported by a Retrieval-Augmented Generation (RAG) knowledge base and a Next.js frontend, all written in TypeScript.

Standardizing Agentic Tool Use with MCP

The Model Context Protocol (MCP) represents a paradigm shift in how Large Language Models interact with external data environments. Historically, enabling an LLM to browse the web, query a specific database, or execute code required custom, hardcoded API wrappers that were brittle and difficult to scale across multiple specialized agents. MCP standardizes the communication layer between AI agents and diverse data sources, allowing for seamless context sharing, tool execution, and resource discovery across distributed networks.

The developer’s deep technical engagement with MCP and CLI-based AI tools is explicitly demonstrated by their activity in the google-gemini/gemini-cli repository. The developer identified, documented, and escalated a critical failure in the GitHub MCP server discovery process. According to the bug report, the connection failed during the discovery phase due to a 'Streamable HTTP error' yielding a 'bad request' status. The system specifically indicated that the 'Authorization header is badly formatted' when attempting a POST request to the endpoint, which entirely prevented the GitHub MCP from connecting successfully. This level of intricate debugging indicates a rigorous, low-level understanding of how MCP servers authenticate, manage headers, and stream data across distributed cyberphysical networks.

Subagent Architecture and the Prevention of Context Rot

Furthermore, the developer actively evaluated and praised the introduction of "Subagents" in the Gemini CLI v0.38.1 release. The Subagent architecture allows a primary LLM process to delegate complex, specialized tasks to isolated expert agents operating within their own unique context windows. This architectural pattern is designed to prevent "context rot"—the degradation of reasoning quality, instruction adherence, and recall that inevitably occurs when a single agent is flooded with too much diverse information within a single session.

The developer noted the efficacy of this release: "The Subagents architecture is a huge win for maintaining session speed, and using Markdown for custom definitions makes it very accessible for power users. Shoutout to @jacob314 for the UX polish—fixing the terminal flicker and adding the input scrollbar makes a world of difference for high-volume work".

When this 13-agent, 14-MCP-server subagent architecture is transposed to a Tokenized Bio Hedge Fund, it becomes the analytical engine of the entire enterprise. Evaluating early-stage biotechnology investments requires parsing heavily obfuscated and mathematically dense datasets: FDA regulatory filings, genomic sequencing data, patent literature, and double-blind clinical trial statistics. A monolithic AI would immediately succumb to context rot under this cognitive load.

Instead, a Subagent architecture allows a primary "Portfolio Manager" agent to orchestrate operations asynchronously. It delegates tasks to deeply specialized subagents: one subagent (operating via an NCBI/PubMed MCP server) analyzes peer-reviewed literature for mechanism-of-action viability; another subagent (via an on-chain analytics MCP server) models the liquidity pools of the target protocol; and a third (via a legal database MCP) assesses patent exclusivity periods. The primary agent synthesizes these isolated, high-fidelity context streams into a unified investment decision.

The Theoretical Engine: The Free Energy Principle and Active Inference

To fully comprehend the theoretical ceiling of the systems built by this developer, it is necessary to analyze the integration of Karl Friston's Active Inference and the Free Energy Principle (FEP). Extensive research materials surrounding the developer's ecosystem highlight a 2021 symposium featuring Professor Karl Friston, organized by the Active Inference Lab. Active Inference provides a mathematically rigorous, physics-based framework for understanding how sentient systems—including artificial multi-agent networks—operate, learn, structure their environments, and survive in highly volatile niches.

The Free Energy Principle as a Variational Principle of Least Action

The Free Energy Principle posits that any self-organizing system that successfully resists the natural entropic tendency toward disorder must do so by minimizing its variational free energy. In the realm of statistical thermodynamics and quantum mechanics, this is recognized broadly as a variational principle of least action. Variational free energy acts as a mathematically tractable upper bound on surprise (or, in statistical terms, the negative log marginal likelihood). By continuously minimizing free energy, a system effectively maximizes the evidence for its own internal generative models of the world.

Friston emphasizes a crucial distinction between the FEP and other variational principles (such as Hamilton's principle of least action or Noether's Theorem) by focusing on the separation of states. The FEP applies specifically when a universe of states is partitioned into the states owned by an agent (internal states), states owned by the environment (external states), and the sensory/active states that mediate the exchange between them. This boundary is mathematically formalized as a Markov blanket. Because the system is separated from the external environment by this Markov blanket, it cannot directly access or measure external states. It can only infer them based on the sensory impressions that penetrate the blanket. Therefore, the internal dynamics of the system necessarily encode Bayesian beliefs about the hidden causes of its sensory data.

Dual-Aspect Information and Information Geometry

This formulation moves beyond standard predictive processing or classical Shannon information theory. Friston differentiates "Information of the first kind" (Shannon information, self-information, surprisal) from "Information of the second kind" (Information Geometry). Shannon information deals merely with the implausibility of a particular event or message. Information geometry, however, posits that the internal state of a computer or neural network can be read as encoding a posterior belief about latent states outside the Markov blanket. This creates a statistical manifold where any movement—any processing of data—necessarily implies Bayesian belief updating. This equips the artificial system with a rudimentary mathematical sentience, where its internal representations possess "aboutness" regarding the external world.

Active Inference and Teleological Planning

Active Inference serves as the "teleological spin-off" of the Free Energy Principle. While passive inference involves merely updating internal beliefs to match incoming sensory data (perception), Active Inference asserts that systems also *act* upon the world to generate sensory data that conforms to their prior beliefs (action). It is the circular causality of engagement between an agent and its environment.

Crucially, when planning behaviors into the future, an agent must optimize an objective functional known as *Expected Free Energy* (EFE). EFE combines two distinct computational imperatives:

1. **Epistemic Value (Information Gain):** The agent actively seeks out observations that resolve uncertainty about the world (a process known as epistemic foraging).
2. **Pragmatic Value (Prior Preferences):** The agent seeks out states that align with its evolutionary, structural, or programmed goals.

Unlike standard Reinforcement Learning (RL), which relies heavily on a state-action value function to maximize explicitly coded rewards, Active Inference unifies both exploration and exploitation under a single, rigorous mathematical imperative: the continuous reduction of uncertainty. By relying on a generative model that rolls out into the future, the agent evaluates the expected free energy of various policy trajectories and selects the path of least action.

Structure Learning and Bayesian Model Selection

A significant challenge in deploying multi-agent AI networks is determining the correct model architecture—often referred to as the structure learning problem. A developer must decide the factorization of the model, the necessary conditionally independent factors, and the appropriate level of granularity to capture latent causes without incurring computational explosion. Friston suggests that the FEP inherently solves this through Bayesian model selection. By utilizing discrete state-space models and message-passing schemes on Forney factor graphs (such as those developed in ForneyLab), systems can naturally prune unnecessary complexity, finding the simplest possible model that accurately explains the data. This adherence to Ockham's principle minimizes the statistical complexity of the system, ensuring rapid, efficient processing.

Mental Action and the Deployment of Precision

A highly advanced concept within Active Inference is the deployment of precision, which functions as a normative theory for "mental action". In hierarchical generative models, an agent must not only infer the state of the world but also infer the reliability (precision) of its own sensory inputs and lower-level beliefs. By optimizing estimates of posterior precision, the system essentially learns what to pay attention to and what to ignore.

Friston points out that this is functionally what transformer networks do via attentional selection mechanisms. However, framing it as mental action within an Active Inference model allows the agent to formulate meta-inferences—predicting its own inference processes. This is directly analogous to the developer's appreciation for the Subagent architecture in Gemini CLI. By delegating tasks to isolated subagents, the primary system effectively executes a mental action, selectively deploying precision to specific, partitioned context windows to resolve ambiguity in a computationally bounded manner, mimicking the cognitive offloading seen in biological brains.

Generalized Synchrony and Multi-Agent Consensus

When scaling from a single agent to an ecosystem of agents, the concept of the "synchronization manifold" becomes paramount. When multiple agents interact—whether dancing in a concert, communicating via social media, or executing trades in a market—they engage in mutual belief updating across their respective Markov blankets. This reciprocal coupling leads to generalized synchrony.

Achieving this synchrony minimizes free energy for all participants because the environment (composed of other agents) becomes mutually predictable. In multi-agent modeling, achieving a shared generative model allows diverse subagents to coordinate seamlessly. Interestingly, Friston notes that the only evolutionarily stable partitioning of populations is often a 50/50 split (e.g., in-groups and out-groups), as any minority group is eventually absorbed by the majority due to the dynamics of marginal likelihood maximization. This mathematical quirk highlights the necessity of carefully architecting consensus mechanisms in multi-agent networks to prevent total cognitive homogenization while maintaining functional synchrony.

Architecting the Tokenized Bio Hedge Fund: A Synthesis

Given the inaccessibility of the exact Tokenized-Bio-Hedge-Fund repository, we must mathematically and architecturally project its structure by fusing the developer's established competencies: the dual-penalty bonding curve, the 13-agent MCP subagent architecture, the stablecoin ecosystem, and the theoretical imperatives of Active Inference.

A hedge fund is fundamentally an engine constructed to minimize financial uncertainty while maximizing asset growth over time. A *Bio* Hedge Fund deals with uniquely volatile, high-uncertainty assets, such as prolonged drug discovery pipelines, synthetic biology patents, and regulatory approval pathways. A *Tokenized* Bio Hedge Fund distributes the risk, verification, and governance of these assets across a decentralized cryptographic network.

The Hedge Fund as a Macroscopic Active Inference Agent

In this architectural projection, the hedge fund itself operates as a macroscopic Active Inference agent bounded by a systemic Markov blanket.

- **External States:** The hidden causes in the world, including global macroeconomic conditions, biotechnology clinical trial outcomes, and shifting FDA regulatory frameworks.
- **Sensory States:** The real-time data streams ingested via the 14 MCP tool servers (e.g., Bloomberg terminals, patent databases, open-source genomic repositories).
- **Internal States:** The generative model of the fund, physically represented by the coordinated consensus state of the 13 specialized AI subagents.
- **Active States:** The execution of capital allocation, token minting/burning, and dynamic, algorithmic adjustments to the bonding curve parameters.

The primary directive of this autonomous fund is to minimize its Expected Free Energy. When faced with an investment decision in a novel biotech protocol, the fund encounters massive uncertainty. Standard algorithmic trading bots, which rely on historical pricing data or simple

momentum indicators, fail in this environment because novel IP lacks historical price action.
An Active Inference agent, however, engages in rigorous *epistemic foraging*. It mathematically calculates which specific piece of data will maximally reduce its uncertainty regarding the latent states of the biotech asset. The primary orchestrating agent delegates this task to a specialized subagent. For example, it instructs an @FDA_Investigator subagent to execute a highly targeted query via an MCP server to extract specific clinical trial phase II toxicity data. This directed foraging satisfies the epistemic value of the Expected Free Energy equation.

Economic Niche Construction via the Dual-Penalty Bonding Curve

As Friston notes, self-organizing systems engage in niche construction—altering their physical or social environment to make it more predictable and thus minimizing future surprise. The Tokenized Bio Hedge Fund engages in profound economic niche construction through its dual-penalty-bonding-curve.
By enforcing market-aware exponential decay and graduated penalties, the fund algorithmically alters the behavior of human traders interacting with its liquidity pools. It penalizes volatile, short-term speculation (which increases systemic free energy and disorder) and financially rewards long-term holding (which decreases systemic free energy and stabilizes the fund's internal generative model). The bonding curve acts as a teleological control mechanism, an active state projecting outward across the Markov blanket to suppress financial entropy in the fund's immediate economic niche.

Subagent Subject Matter Experts vs. Prediction Matter Experts

Friston's symposium touches upon an intriguing distinction between "Subject Matter Experts" and "Prediction Matter Experts". In highly volatile environments, agents tend to become prediction experts, assigning high precision to recent evidence to rapidly adapt to a capricious world. In stable, predictable environments, agents minimize the complexity of their generative models and become subject matter experts, highly tuned to a specific, narrow domain.
The 13-agent architecture of the TRAVEL-Planner, transposed to the hedge fund, likely utilizes a mix of both. Subagents tasked with analyzing stable, well-documented domains (e.g., existing patent law) operate as Subject Matter Experts with rigid, low-complexity generative models. Conversely, subagents tasked with parsing real-time market sentiment or live clinical trial leaks operate as Prediction Matter Experts, optimized for high-volatility epistemic foraging. The primary agent synthesizes these distinct cognitive styles to maintain a balanced, comprehensive view of the market.

Implications for Decentralized Science (DeSci) and Financial Modeling

The synthesis of multi-agent AI, robust cryptographic bonding curves, and Active Inference represents a quantum leap in the rapidly expanding field of Decentralized Science (DeSci). Traditional biotech hedge funds rely on teams of human analysts who are inherently subject to cognitive biases, bounded rationality, inter-group political friction, and limited context windows. An autonomous, tokenized fund modeled on Bayesian mechanics transcends these biological limitations. It does not merely react to the market; it actively infers the hidden causes of market movements and biological research outcomes, adjusting its internal generative model to

minimize surprise. Furthermore, by democratizing access to this fund via tokenization, retail participants and independent researchers can gain exposure to institutional-grade biotechnology investments. Importantly, these participants are protected from predatory DeFi mechanics—such as the zero-duration extraction of flash loans—by the sophisticated temporal friction of the dual-penalty bonding curve.
The integration of Friston's theoretical principles suggests that the ultimate future of algorithmic trading is not traditional deep learning, which passively maps massive datasets to outputs via black-box weights. Instead, the future belongs to generative modeling that supports counterfactual planning. A financial entity that can query its own model, ask "What would happen if I intervened in this specific biotech market?", and mathematically bound the uncertainty of that counterfactual, operates at a level of sophistication previously relegated to theoretical neuroscience and advanced physics.

Conclusion

The extensive digital footprint and technical engagements of coad1024-cmd outline the architectural blueprint for a highly resilient, autonomous financial entity. While the specific codebase for the Tokenized-Bio-Hedge-Fund repository remains hidden from public view, the surrounding ecosystem provides ample evidence of its systemic design. This ecosystem comprises the anti-fragile dual-penalty-bonding-curve, the intricate TRAVEL-Planner multi-agent MCP framework, the verifiable asset generation of skillproof-backend, and foundational research into stablecoin economic base layers.
By applying the rigorous mathematical principles of Active Inference and the Free Energy Principle, this architecture transcends standard smart-contract automation. It represents a teleological, self-organizing financial agent capable of directed epistemic foraging, Bayesian structure learning, precision deployment via mental action, and dynamic niche construction. As decentralized technologies and large language models continue to mature, frameworks that successfully marry the predictive power of multi-agent subagent architectures with the systemic resilience of advanced tokenomics will likely dominate the future of asset management, particularly in volatile, data-dense sectors such as global biotechnology.

Works cited

1. Dzuname coad1024-cmd - GitHub, https://github.com/coad1024-cmd 2. GitHub · Where software is built, https://github.com/orgs/causalincentives/followers 3. skillproof-backend/GEMINI.md at main - GitHub, https://github.com/coad1024-cmd/skillproof-backend/blob/main/GEMINI.md 4. Views · coad1024-cmd/dual-penalty-bonding-curve · GitHub, https://github.com/coad1024-cmd/dual-penalty-bonding-curve/issues/views 5. Discovery Error: GitHub Connection Failed due to Bad Authorization Header #14688, https://github.com/google-gemini/gemini-cli/issues/14688 6. Gemini CLI v0.38.1: Subagents are here #25562 - GitHub, https://github.com/google-gemini/gemini-cli/discussions/25562